# Tokens-per-Parameter Coverage Is Critical for Robust LLM Scaling Law Extrapolation

Joshua Shay Kricheli* Alexander Lawrence Reid* Soumajyoti Sarkar
Venkata Gandikota Paulo Shakarian

Syracuse University: Leibniz Lab, Institute for Artificial Intelligence Amazon AGI Foundations†
*Equal contribution. †This work does not relate to the author's position at Amazon.

NeurIPS 2026, Sydney





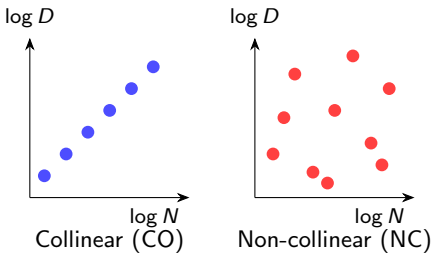







# Collinear designs

- Scaling laws are fitted on small runs and used to predict large ones.
- Many studies train at a fixed tokens-per-parameter ratio $k = D/N$: a ray.
- On a ray both power-law terms move together:
  $L(N, kN) = AN^{-\alpha} + Bk^{-\beta}N^{-\beta} + E.$
- Only the exponent gap $\varepsilon = |\alpha - \beta|$ separates $A$ from $B$.



# Theory: when the scale pair is identified

- One ray: $\kappa_{A,B} = \Theta(\varepsilon^{-2})$, and the interval on $A$ is $\Omega(\varepsilon^{-1})$ times the one on the combined $\psi = A + Bk^{-\alpha}$.
- $K$ rays $k_1 < \cdots < k_K$: to leading order $\kappa_{A,B} \le \kappa_{\mathrm{target}}$ whenever

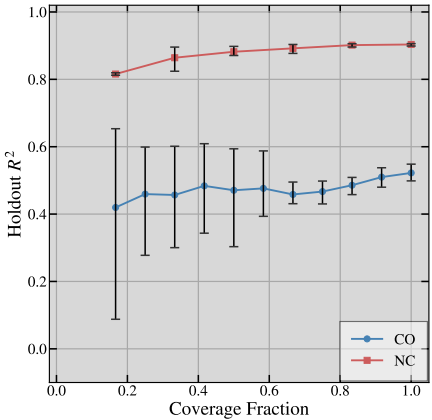
$$\underbrace{\frac{1}{K}\sum_{\ell=1}^{K} k_\ell^{-2\beta_{\mathrm{eff}}} - \Big(\frac{1}{K}\sum_{\ell=1}^{K} k_\ell^{-\beta_{\mathrm{eff}}}\Big)^2}_{V_K} \;\ge\; \underbrace{\frac{\big(K + \sum_\ell k_\ell^{-2\beta_{\mathrm{eff}}}\big)^2}{K^2\,\kappa_{\mathrm{target}}}}_{\tau_K},$$

- and $\kappa_{A,B} \le \kappa_{\mathrm{target}}$ iff $V_K \ge \tau_K/(1 + 1/\kappa_{\mathrm{target}})^2$.
- One ray never passes; two rays pass once $k_2/k_1 \ge R_{\min}(k_1)$, a check made before training.

# Results: non-collinear designs extrapolate

| Comparison (NC vs. CO) | NC wins |
|---|---|
| Full fan, FP32 (1,460 of 1,500) | 97.3% |
| Full fan, BF16 (296 of 300) | 98.7% |
| Budget-matched, Chinchilla | 68.5% |
| Budget-matched, Droppo-Elibol | 57.4% |
| Budget-matched, Kaplan | 80.6% |

- About 1,900 models, 5.04M-76.5M parameters, five corpora, four laws.
- Budget-matched rows: CO subsets below the threshold.




Holdout $R^2$ against TPP coverage (Kaplan, RedPajama).

# Takeaway and design recipe

- Report coefficients only from designs that span TPP ratios.
- Recipe: pick $k_1$, read $R_{\min}(k_1)$, train two ratios at the ends of that spread, split runs evenly, and check the exact $\kappa_{A,B}$ before training.
- Code, data and checkpoints are public.

`lab-v2.github.io/tpp-coverage`

